



RESEARCH WITH INKHAVEN

Grounding Your Book in Fact

Researching Fiction and Non-Fiction with Inkhaven's Research Assistant

Vladimir Ulogov

2026 · examples assume Inkhaven 1.5.9 or newer

Contents

Who this book is for	1
What you will be able to do	2
What you will need	2
The one idea to carry with you	3
How to read this book	4
Part I — Grounding Your Book	5
1 — Why Ground Your Book	6
Two kinds of claim	7
Not all grounding is equal	8
Where your facts live	9
Provenance: the fact that remembers	10
2 — The Research Assistant	12
Two panes, one conversation	13
Two ways to speak	13
The arc of a project	14
Nothing happens to your manuscript by accident	15
3 — Your First Fact	16
Ask	17
Turn the answer into a candidate fact	18
The gate: you get the last word	18
What you just kept	19
The habit under the habit	19
About the Author	21
Why the Research Assistant exists	22
A work of love	22
A note on cooperation	23
Where to find more	23

Before You Begin

This is a book about getting the facts right. Not the invented facts — the ones you make up freely, which are yours to command — but the **borrowed** ones: the year a bridge was built, the population of a city, the distance a rider could cover in a day, the claim in your argument that a reader could look up and check. Every book leans on facts it did not invent, and a book that gets them wrong loses the reader's trust in the ones it did.

You will do this inside **Inkhaven** — a writing tool with a built-in **Research Assistant**: a quiet workspace where you ask questions, gather answers from trustworthy places, check them against each other, and keep only what survives. By the end of this book you will be able to take a manuscript from “I think that’s about right” to a knowledge base where every fact you kept remembers **where it came from** — and to compose from that knowledge base directly into your writing.

Who this book is for

This guide assumes **no prior knowledge** — of research, or of Inkhaven.

You do not need to be a researcher. You do not need to know what a citation manager is, or how a database works, or the difference between a primary and a secondary source. Every idea is introduced the first time it is needed, in a marked box like this one:

TERM **Fact**

In this book, a **fact** is a specific, checkable claim you want your writing to rest on — a date, a number, a name, a cause. It is the opposite of a thing you invented. The whole craft here is telling the two apart, and grounding the first kind.

You do not need to be a novelist, either. The Research Assistant serves two kinds of author equally, and this book keeps both in view the whole way through. Where

a task looks different depending on what you write, you will see it split into two tracks:

FICTION

You are writing a novel set during the building of a Roman aqueduct. You need the engineering to feel **true** — real distances, plausible dates — without stopping the story to fact-hunt.

NON-FICTION

You are writing a history of Roman water supply. You need every claim **sourced** — a real citation behind each figure — so a reviewer can follow your evidence.

Same tool, same workflow, two destinations: the novelist wants the world to feel solid; the non-fiction author wants the argument to be defensible. Both are the same act — grounding a claim — and both are what the Research Assistant is for.

What you will be able to do

By the last page you will be able to:

- ask a question and get an answer **grounded on facts you already trust**, not a confident guess;
- pull facts from authoritative places — a structured knowledge base, real places, scholarly papers, public-domain books, the open web — each one remembering its source;
- cross-check a shaky claim against several independent sources before you commit it, and let the tool argue **against** a fact to see if it survives;
- keep your knowledge base honest over time — re-grounding old guesses, flagging facts that may have gone stale;
- and compose **out** of the corpus you built: a cited synthesis, an outline, a real bibliography — straight into the book you are writing.

What you will need

Inkhaven, installed and open on a project — a folder with a manuscript in it. That is all. Nothing in this book requires an account, a subscription, or an internet connection you have to pay for. A few features reach out to free public services (a maps database, a scholarly index, a library of public-domain books); each one is


optional, degrades gracefully when you are offline, and is clearly marked where it appears.

ON THE AI

The Research Assistant uses a language model to help you **find** and **phrase** answers — but it never decides what is true for you, and it never edits your prose. Every fact crosses a confirmation step before it is kept, and every fact records where it came from. The model is a research partner with a short leash, not an oracle. If you have no model configured, most of the workflow still works; the parts that need one say so.

The one idea to carry with you

If you remember nothing else from this introduction, remember this picture. It is the spine of the whole book:



The trust ladder — every fact you keep carries a rung, recorded as its provenance.

Not all facts are equally trustworthy, and the Research Assistant never pretends otherwise. A number you **computed** is firmer than one a structured database gives you; that is firmer than a scholarly paper; that is firmer than a web page; and all of them are firmer than a model's unaided guess. Every fact you keep is stamped with its rung on this ladder — its **provenance** — so that later, you (or a reader, or a reviewer) can see at a glance how much weight it can bear.

Everything else in this book is a way of climbing that ladder: getting a claim from a guess at the bottom to something cited and checked near the top.

How to read this book

Read it front to back the first time. Part I builds the foundation — why grounding matters, the workspace you will live in, and your very first grounded fact — and everything after it assumes those three chapters. Later parts are more independent: once you have the basics, you can jump to the source you care about (real places, scholarly papers, the web) or the technique you need (cross-checking, composing a bibliography).

Commands you type look like `/fact` or `/geonames`. Keys you press look like `Ctrl+S`. When a chord opens something on screen, the book tells you what you would see and why it matters — but it shows you the **shape** of the work with diagrams rather than screenshots, because screenshots age badly and the ideas underneath them do not.

Turn the page, and let us talk about why any of this matters.

WHAT YOU LEARNED

- A **fact** here is a borrowed, checkable claim — the opposite of what you invented — and grounding it is the whole craft.
- The book serves **fiction and non-fiction** authors equally, and assumes no prior knowledge of research or Inkhaven.
- Facts differ in trustworthiness; the **trust ladder** ranks them, and every kept fact records its rung (its **provenance**).
- The AI helps you find and phrase, but never decides truth and never edits your prose — every fact crosses a confirmation step.

GROUNDING YOUR BOOK IN FACT

PART I

Grounding Your Book

CHAPTER 1

1

Why Ground Your Book

A reader gives you enormous latitude. They will follow you into an invented world, accept a magic system, believe in characters who never lived, forgive a coincidence that would never happen. What they will not forgive is being made to feel foolish — and nothing does that faster than a fact they **know** to be wrong.

The moment a reader catches you saying the aqueduct carried water uphill, or that a medieval army marched two hundred miles in a day, or that a chemical does something it cannot do, a small door closes. Not because the error is large, but because it reveals that you did not check — and if you did not check **that**, what else did you not check? The invented facts and the borrowed facts sit side by side on the page, and the reader cannot always tell which is which. So they trust all of them,

or none. The borrowed facts are load-bearing precisely because a reader **can** look them up.

Two kinds of claim

Every sentence you write that makes a claim about the world is one of two kinds.

Some claims are **yours**. The name of your protagonist, the shape of your invented coastline, the rule that your magic costs the caster a memory — these you command absolutely. There is no external truth to get wrong. We will come back to these in a later chapter, because they deserve protection **from** fact-checking, not by it.

Other claims are **borrowed**. They rest on a world that exists independently of your book — history, geography, science, the record of what actually happened. You did not invent the speed of a horse or the year Rome fell, and a reader can check them. These are the claims this book is about.

TERM Grounding

Grounding a claim means connecting it to something outside your own head that supports it — a source, a computation, a cross-checked agreement between independent references. A grounded claim can answer the question “how do you know?” An ungrounded one can only answer “it felt right.”

The first discipline of research, then, is simply **telling the two apart**. A surprising amount of trouble comes from treating a borrowed claim as if it were yours — writing “the journey took three days” because it felt right, when the distance and the terrain would have made it eight. The Research Assistant cannot tell you which claims to borrow. But once you know a claim is borrowed, it gives you everything you need to ground it.

FICTION


In a novel, grounding is **invisible craft**. The reader never sees your sources; they feel the result — a world that holds together, that never trips them. You ground so the seams don’t show.

NON-FICTION

In non-fiction, grounding is **the argument itself**. The reader is meant to see your sources; a claim without one is an assertion, and assertions don’t persuade. You ground so the case can be followed.

Not all grounding is equal

Here is the idea from the introduction again, because everything hangs on it. When you ground a claim, **what you ground it on matters**. A source can be more or less trustworthy, and the Research Assistant is honest about the difference at every step. It arranges the ways of knowing into a ladder:



The trust ladder — every fact you keep carries a rung, recorded as its provenance.

Read it from the bottom up.

Model — an educated guess

At the bottom sits the language model answering from memory. It is often right, frequently fluent, and occasionally confidently wrong — and it cannot tell you which. A model's unaided answer is a **starting point**, never a resting place. Useful for finding your question; never sufficient for keeping the answer.

Web — a cited page

A step up: an answer drawn from actual web pages, with the pages named. Now there is something to check. The Assistant fact-checks a web-grounded claim before it lets you keep it, and records the page it came from.

Documents — sources you imported

Higher still: your own trusted materials — a PDF, a research folder, a book — brought into the corpus so answers can be drawn from **them** specifically, and cited back to the page.

Scholarly — a real paper

A named work in the scholarly record (a journal article, a preprint), with an identifier a reviewer can resolve. The claim now rests on the literature.

Structured — cited by id

A structured knowledge base — geographic gazetteers, curated fact databases — where a claim comes as a discrete, identified datum rather than a paragraph of prose. Not “an article that mentions Rome,” but “Rome, id Q220, founded –753.”

Computed · Simulation — deterministic

At the top, facts you did not look up but **derived** — a distance from coordinates, a travel time from a speed, a fact your own world-simulation produced. These are the firmest of all, because anyone can re-run the computation and get the same answer.

WHY A LADDER AND NOT A WALL

The ladder does not mean “only ever use the top rung.” A novelist grounding the feel of a city may be perfectly served by a web page; a historian needs the scholarly rung or higher. The point is not to forbid the low rungs — it is to **know which rung you are standing on**, and to record it, so nothing pretends to be firmer than it is.

Where your facts live

Grounded claims have to go somewhere. In Inkhaven they live in a small set of places, kept separate on purpose:

- **Facts** — the book you trust. A claim only lands here after it has been grounded and you have confirmed it. This is your ground truth.
- **Notes** — the book you are still unsure about. A promising claim that has not earned the Facts book yet; a lead to chase later.
- **Sources** — your bibliography. The citable works behind your facts — papers, books, imported references — accumulated as you research, ready to become a real reference list.

TERM Corpus

Your **corpus** is the whole body of research material you have gathered for a project — the Facts, the Notes, the Sources, and the imported documents behind them. It grows as you work and, crucially, you can search and compose from it. It is the difference between “I read a lot for this book” and “I can find and cite exactly what I read.”

The separation matters. Keeping speculative Notes out of your trusted Facts is what lets you **rely** on the Facts — when you pull a fact into your manuscript, you are pulling from a book you vetted, not from a pile where the checked and the unchecked are mixed together.

Provenance: the fact that remembers

The single most important habit this book will build is small and almost invisible: **every fact you keep remembers where it came from.**

TERM Provenance

Provenance is the recorded origin of a fact — which rung of the trust ladder it came from, and the specific source (a page id, a paper, a URL, a computation). It travels with the fact silently, and you can call it up any time to ask a fact “how do you know?”

Provenance is what turns a knowledge base from a heap into an instrument. Six months from now, when a copy-editor queries a date, you will not have to reconstruct where you found it — the fact will tell you. When you assemble a bibliography, it will build itself from the provenance you accumulated. And when a fact is only a model’s guess, its provenance says so plainly, so you know it still needs work.

You will rarely type the word “provenance” again. But it is running underneath every command in this book, and it is the reason the Research Assistant can be trusted to be honest about what it knows.

WHAT YOU LEARNED

- Borrowed (checkable) claims are load-bearing: one caught error makes a reader doubt the rest, invented and borrowed alike.
- The first discipline is telling **invented** claims (yours to command) from **borrowed** ones (which must be grounded).
- The **trust ladder** ranks ways of knowing from a model's guess up to a deterministic computation; the goal is to know — and record — which rung you stand on.
- Facts (trusted), Notes (speculative), and Sources (citations) are kept separate so you can rely on the Facts.
- **Provenance** — every kept fact remembering its origin — is the quiet habit that makes the whole corpus trustworthy.

CHAPTER 2

2

The Research Assistant

The Research Assistant is a separate screen inside Inkhaven — a room you step into when you want to research, and step out of to write. It is deliberately quiet: two panes, a place to type, and a status line. Everything you do here either **gathers** a fact, **checks** one, or **composes** from the ones you have kept.

This chapter is a tour of the room and the way you move through it. There is no command to memorise yet; the point is to understand the shape of the work so the commands in the next chapter feel like the natural thing to type.

Two panes, one conversation

Picture the screen split in two.

On the left is your **Facts tree** — the growing outline of everything you have kept for this project: your Facts, organised the way a book is, in chapters and sections you can fold and unfold. When you are new to a project this pane is nearly empty. It fills as you work, and by the end of a book it is the map of everything you learned.

On the right is the **conversation** — where you ask questions and read answers. You type at the bottom; the exchange scrolls above. This is where the language model lives, where search results arrive, where a fact-check verdict appears. It is a chat, but a chat with a purpose: everything said here is a candidate for the tree on the left.

The whole workflow is the traffic between these two panes: a question on the right becomes, after checking and your confirmation, a fact on the left.

TERM Thread

A **thread** is one research conversation, saved. You might keep one thread per topic — “the aqueduct,” “1920s Vienna,” “cell biology for chapter 9” — so each line of inquiry has its own history you can return to. Threads keep unrelated research from tangling together, and they persist between sessions.

Two ways to speak

You say things to the Assistant in one of two registers, and learning the difference is most of learning the tool.

The first is **plain language**. You type a question the way you would ask a knowledgeable friend — “how far could a Roman legion march in a day?” — and you get an answer, grounded on whatever you have already gathered. This is for **thinking**: exploring, orienting, working out what you actually need to know.


The second is a **command**. A command starts with a slash — `/fact`, `/geonames`, `/triangulate` — and it does something precise and repeatable: keep this claim as a fact, look this place up in a gazetteer, cross-check this against three sources. Commands are for **acting**: turning a loose exchange into something durable.

YOU DON'T HAVE TO MEMORISE THE COMMANDS

Type a single `/` and a **command palette** opens — a searchable list of every command with a one-line description. You can browse, filter by typing, and pick one without ever having learned its name. Throughout this book, when a command is introduced, that palette is where you would rediscover it later. Learn the workflow; let the palette remember the spelling.

The arc of a project

Zoom out from a single question to a whole book, and the work has a natural shape. It is a loop, and the Research Assistant is built to carry you around it:



The corpus you build is not a dead pile of notes — it is checked, kept fresh, and composed back out.

You **acquire** — ask, search, import — pulling raw material in. You **cross-check** — hold a claim up against independent sources, even argue against it — so that only what survives is kept. You **maintain** — because a knowledge base decays; guesses can be re-grounded on real sources later, and old facts can be flagged for a second look. And you **compose** — turning the corpus back into output: a synthesis, an outline, a bibliography that feeds the very book you are writing.

Most of this book is a slow walk around that loop. Part I gets you onto it — one question, one checked fact. Later parts deepen each quarter of the circle.

FICTION

A novelist tends to live in **acquire** and **compose**: gather enough real texture to make the world convincing, then let it quietly inform the prose. Cross-checking matters most for the load-bearing facts a sharp reader would catch.

NON-FICTION

A non-fiction author lives in **cross-check** and **maintain**: every claim must survive scrutiny and stay current, and the bibliography that **compose** produces is often part of the deliverable itself.

Nothing happens to your manuscript by accident

One promise, made early because it shapes how freely you can experiment: the Research Assistant never touches your manuscript on its own, and never edits your prose. It writes only to your research books — Facts, Notes, Sources — and only the deliberate acts you take (keeping a fact, recording a source) change anything on disk. You can ask anything, explore any tangent, take any command back out for a test drive, and know that your book on the other screen is exactly as you left it.

That safety is what lets the next chapter be hands-on. You are about to gather your first fact — and there is nothing you can do in the process that you cannot undo.

WHAT YOU LEARNED

- The Research Assistant is a separate, quiet screen: a **Facts tree** on the left, a **conversation** on the right, and the traffic between them is the workflow.
- **Threads** keep separate lines of inquiry apart and persist between sessions.
- You speak in **plain language** to think and explore, or in **slash commands** to act; a / palette means you never have to memorise a command.
- A project moves around a loop — **acquire** → **cross-check** → **maintain** → **compose** — and the tool is built to carry you around it.
- The Assistant writes only to your research books and never edits your prose, so you can explore without risk.

CHAPTER 3

3

Your First Fact

Enough groundwork. Let us put one real fact into your book, from question to kept claim, and watch what the Research Assistant does at each step. The whole thing takes about a minute; understanding it takes a chapter, because every stage is teaching you something you will lean on for the rest of the project.

Here is the shape of what we are about to do:



Every claim earns its place: nothing is written to your Facts without your say-so.

Ask a question. Ground the answer on what you already trust. Cross a confirmation step — the **gate** — where you get the final say. Keep the result as a fact that remembers where it came from. Four steps, and you will do them thousands of times.

Ask

Open the Research Assistant, and in the conversation pane type a plain question — whatever your book actually needs to know. Not a command, not a keyword; a real question.

FICTION

How far could a Roman legion march in a single day on a good road?

You need this because your protagonist has to reach the frontier “in time,” and **in time** has to mean something.

NON-FICTION

What was the average daily marching distance of a Roman legion, and what is the evidence for it?

You ask for the **evidence** too, because your reader will expect a citation, not just a number.

The Assistant answers in prose. Behind the scenes it did something important before replying: it looked through the facts you have **already** kept for this project and offered them to the model as context, so the answer is consistent with your own established world — not a generic answer that might contradict what you decided three chapters ago.

TERM Grounding on your corpus

Before answering, the Assistant **retrieves** the most relevant facts you have already kept and hands them to the model along with your question. This is why its answers get better as your corpus grows: it is not guessing in a vacuum, it is reasoning over what **you** have established. Early on, with an empty corpus, it leans on the model’s own knowledge — which is exactly why the next steps matter.

Turn the answer into a candidate fact

Suppose the answer says a legion covered roughly twenty Roman miles a day, more under forced march. That is useful — but right now it is just words in a chat, sitting on the bottom rung of the ladder: a model’s guess. To keep it, you turn it into a candidate fact with a command:

```
/fact a Roman legion marched about 20 Roman miles (30 km) per day
```

You do not have to retype the whole answer; you can point `/fact` at the claim you want and let it distill a clean, single statement. What you are doing is **promoting** a line from the disposable conversation into something the book will remember.

FACT OR NOTE?

If the claim is solid enough to build on, `/fact` puts it in your trusted Facts book. If it is a promising lead you are not sure about yet, its sibling `/note` puts it in Notes instead — the same gesture, a different shelf. You can promote a Note to a Fact later, once it has earned it. When in doubt, take a Note; the Facts book is worth keeping clean.

The gate: you get the last word

Here is the step that makes the whole tool trustworthy. `/fact` does not silently write to your book. It opens a **confirmation** — the gate — showing you exactly what is about to be kept: the fact’s wording (which you can edit), where it will be filed, and where it says it came from.

TERM The confirmation gate

The **gate** is the mandatory pause between “the Assistant proposed a fact” and “the fact is in your book.” Nothing reaches your Facts without passing it. You can edit the wording, accept it with `Ctrl+S`, or discard it entirely. The Assistant proposes; you dispose.

For a plain model-derived claim like this one, the gate may also offer to **check** the fact before you commit — and for facts drawn from the web or from structured

sources, later chapters show the gate doing real cross-checking here. For now the important thing is the principle: **a fact becomes yours only when you say so**. If the wording is off, fix it in the gate. If you have changed your mind, discard it and nothing is written. If it is right, confirm — and it lands on the tree to your left.

What you just kept

Look at the new fact in the Facts tree. It is not just a sentence. Attached to it, quietly, is its **provenance** — a small record that this fact came from the model (the bottom rung), together with the question that produced it. You can call that record up any time.

This matters more than it seems. That provenance is honest: it says, in effect, “this is currently only a guess.” Later chapters will show you how to **climb the ladder** with this exact fact — re-grounding it on a real source with a single command, at which point its provenance updates to say so, and the fact grows firmer without you rewriting a word of it.

FICTION

Your legion’s march is now a kept fact your story can lean on — and one you have flagged (via its provenance) as still needing a real source before you fully trust the timeline that depends on it.

NON-FICTION

Your marching-distance claim is kept, but its provenance says **model** — which for your purposes is not good enough yet. You now know precisely which of your facts still need a citation, because each one tells you.

The habit under the habit

That is the entire core loop, and it never gets more complicated than this: ask, **/fact**, confirm. Everything else in this book makes the **middle** richer — better places to draw the answer from, real cross-checking at the gate, ways to keep the corpus healthy — but the gesture stays the same. Ask a question; keep what survives; let it remember where it came from.

In the next part we start climbing. Instead of grounding on the model’s memory, you will draw facts from authoritative sources — structured knowledge bases, real

places, scholarly papers, whole public-domain books — so that the facts you keep start their life much higher on the ladder.

WHAT YOU LEARNED

- The core loop is **ask** → **/fact** → **confirm**, and it never gets harder than that.
- A plain question is answered **grounded on the facts you already kept**, so answers improve as your corpus grows.
- **/fact** keeps a trusted claim; **/note** keeps a speculative one — when in doubt, take a Note and promote it later.
- The **gate** is a mandatory pause: you edit, confirm, or discard, and nothing reaches your Facts without your say-so.
- Every kept fact records its **provenance** — even “just a model’s guess” — which is what lets you climb the ladder later without rewriting anything.

AFTERWORD

About the author



Vladimir Ulogov.

Vladimir Ulogov has spent decades building infrastructure for distributed systems — the kind of software that watches other software. Early in his career he worked on monitoring and telemetry platforms; later years took him into federated observability, telemetry buses, and the architecture of systems that have to make sense of millions of data points without losing the thread.

Observability, in the end, is a discipline of **grounding** — of never reporting a number the system cannot back up, of knowing the provenance of every signal. It is not an accident that a tool he built for writers carries the same instinct into the library.

What makes him slightly unusual in his corner of the industry is a tendency to write his own tools — not in the sense of small utilities, but

in the sense of programming languages. The Bund language (its compiler, its VM, its document store, its parser) lives in a long series of Rust crates on crates.io. `rust_dynamic`, `rust_multistack`, `rust_multistackvm`, `bundcore`, `zbus_universal_data_gateway` — each is a building block that exists because the off-the-shelf options didn't fit the shape of the work.

Why the Research Assistant exists

Inkhaven is Vladimir's personal reflection on how a literary tool can help the people who write books. The Research Assistant is one answer to a specific frustration: that the tools writers use to **gather** facts and the tools they use to **write** live in different worlds, and the seam between them is where errors and lost citations breed.

A novelist keeps a browser open in one window and a manuscript in another, and the fact they looked up on Tuesday is gone by Thursday. A non-fiction author tracks citations in one program, prose in a second, and prays the two stay in sync. Neither tool remembers **where a fact came from** once it lands in the draft. The Research Assistant was built to close that gap: to make research a room inside the writing tool, where every fact you keep remembers its own provenance and can be composed straight back into the page.

A work of love

Inkhaven is open source. The licence is permissive — you can read it, fork it, study it, modify it, and pass it on. Strictly speaking, the licence also lets you sell it. The author would, gently but firmly, disagree with you doing that. Inkhaven was not designed as a *for sale* project. It is a work of love made for the authors who can least afford to pay for software, and turning it into a commercial product would betray the reason it exists. So please — don't.

It carries no analytics, no telemetry, no upsell. The binary will never phone home; the project will never have an “Enterprise” tier you have to escape from.

This was a deliberate choice. There are excellent commercial tools for research and writing. They cost money — which is fine for many writers, and a barrier for many more. Inkhaven exists for the second group: for the graduate student writing a dissertation on a battered laptop, for the novelist who shouldn't have to pick between rent and software, for the engineer drafting in the same terminal where they already write code, for anyone who would benefit from a tool that respects their work without asking for a credit-card number.

A note on cooperation

Vladimir believes firmly in the human capacity for mutual help — that we make better work, and live better lives, when we share what we know and what we build. Open source is one of the most concrete expressions of cooperation our era has produced: code read, improved, and passed forward without payment, without permission, by people who will never meet.

If Inkhaven helps you finish your book — that is enough. If it gives you a chord pattern you adapt into your own tool — that’s a gift back to the larger project of making software writers can love. As a society, we achieve the greatest things when we help each other rather than compete with each other.

Where to find more

GitHub	@vulogov — the source for Inkhaven, Bund, and the dozen-plus Rust crates that carry the infrastructure. Issues and PRs welcome.
LinkedIn	/in/vladimirulogov — posts on observability, the occasional long-form essay.
YouTube	@vulogov — talks and walkthroughs from the conference trail.
X / Twitter	@vladimir_ulogov

If a fact you grounded ever surprises you, or a command trips you up and you can't find the answer in this book — open an issue on GitHub. The author reads them.