



WORKING WITH INKHAVEN

Theology and Philosophy with Inkhaven

*Researching, Writing, and Composing a Work of Argument — A Worked
Example*

Vladimir Ulogov

2026 · examples assume Inkhaven 1.6.20 or newer

Contents

Introduction	2
The question we will work	3
The sources will be real, and public-domain	3
The loop this book walks	4
Part I — Framing the Question	6
1 — Framing the question	7
Set the genre that stops demanding proof	8
State the question as a thesis you could be wrong about	8
Sharpen it with the Dialectician	9
Part II — Gathering the Sources	12
2 — Gathering the primary sources	13
Kant, from the public domain	14
The scripture adapters	14
Native tongue, native text	16
The same discipline, other doors	16
3 — The corpus and its provenance	18
Facts and notes — two shelves, two confidences	19
Declare your axioms	19
See where everything came from	20
Part III — Interrogating and Reading	22
4 — Interrogating the corpus with SCHOLAR	23
Relate a claim to the corpus	24
Find where the sources contradict	25
Find where they converge	26
Question the whole with the Dialectician	26
The persistent report	27
5 — Reading the draft	29
Confront the paragraph against the corpus	30
Read for theological weight — the Inner Theologian	31
Read the argument for rigor	32
Press the argument — the Socratic personas	33
Part IV — Composing and Producing	35
6 — Composing with loci	36
Cite a locus	37
Build the Index Locorum	38

Validate the loci, and let them canonicalize	39
Fold it into the finished book	40
7 — Revising the prose	42
Snapshot before you revise	43
Grammar-check the passage	44
Read for craft — the Inner Editor	44
The order of the pass	45
8 — Producing the essay	47
The front matter of an essay	48
Render	48
The essay this produced	49
Part V — Closing	51
9 — The discipline, in one page	52
What the tools actually did	52
Where to go next	53
About the Author	56
Why a tool for the unmeasurable	57
A work of love	57
A note on cooperation	58
Where to find more	58

INTRODUCTION

One question, from blank project to finished essay

Theology and philosophy argue about things measurement cannot settle. You are not reporting a fact the reader can check against the world; you are advancing a *position* and asking that it be granted a hearing. That makes the work's obligations unusual — your ground is internal coherence and the tradition you answer to, not a dataset — and it makes the tooling unusual too. A naïve fact-checker pointed at “the soul is immortal” will flag an unsupported claim, which is not a defect in the argument but a category error in the reader. This book is about doing the work with tools that have been recalibrated to the standard the work is actually held to.

It is a *process* book, and it takes an unusual vow: it does not talk about theology-and-philosophy writing in the abstract. It picks one real question and researches and writes it end to end, in front of you, using nothing but Inkhaven. Every stage — framing the project, gathering the primary sources, interrogating them for tension and agreement, drafting against what was found, confronting the draft with the corpus, and producing a finished essay with a bibliography and an index of cited passages — is shown on that one question. The essay the process produces is not hypothetical either: it is the companion volume that sits beside this one on the shelf.

The question we will work

The question is a genuine one in comparative philosophical theology, and it has the shape this track rewards — a real family resemblance between two traditions, running alongside a real disagreement:

ASK YOURSELF

Kant's transcendental idealism and the Latter-day Saint doctrine of *eternal progression* both refuse a single, instantaneous, static salvation — both make perfection a trajectory rather than a state. Where does that resemblance hold, and where does it break? Do the two ever say the same thing, or only rhyme?

We chose it deliberately. It reaches across a philosopher and two scriptural traditions; it turns on what specific passages actually say, so provenance matters; and its whole interest lies in a *graded* relation — agreement here, tension there, flat contradiction elsewhere — which is exactly what the tools in Part III were built to surface. A question with no tension would waste them; a question with only contradiction would be a polemic, not a study.

The sources will be real, and public-domain


Nothing in this book is invented data. Kant comes from Project Gutenberg; the Bible and the Book of Mormon come through Inkhaven's scripture adapters, which draw only on public-domain texts. This is a discipline, not a convenience: a theological argument that recalls what a source says from memory is unreliable and against the whole grounding ethos. We ingest the real texts, retrieve from them with their provenance intact, and let the AI judge only passages actually in front of it — never its own recollection.

INSIGHT

The engine of this track is not a machine that *knows* theology. It is a machine that retrieves what your sources actually say, tests your claims against those passages, and points you to where the sources agree, strain, or contradict. The judgement stays yours. The tools make the judgement *checkable* – and on a track where correctness is off the table, checkable is the closest thing to credible there is.

The loop this book walks

The work is a loop, and this book is its stations in order:



The work is a loop, not a line: you frame the question, gather the sources, interrogate them for tension and agreement, draft against what you found, confront the draft with the corpus, and produce the finished argument — returning to gather more wherever the confront exposes a gap.

We **frame** the question and set the genre that stops the readers demanding proof (Chapter 1). We **gather** the primary sources — Kant, the Bible, the Book of Mormon — as a provenance-tagged corpus (Chapters 2–3). We **interrogate** that corpus with SCHOLAR, the relation engine: where does a claim relate to the sources, where do the sources contradict, where do they converge, and what questions must the whole hold up under (Chapter 4)? We **read** the growing draft — confronting each paragraph against the corpus and weighing it with the Inner Theologian, the reader this track exists for (Chapter 5). We compose with **cited loci** — `@bible[John 3:16]`, `@kant[A51/B75]` — and gather them into an Index Locorum (Chapter 6). We **revise** the prose with the craft reader, the grammar check, and the snapshot beneath a bold rewrite (Chapter 7). And we **produce** the finished essay (Chapter 8).

NOTE

This book assumes you have met Inkhaven's desk before — the chapter tree, the editor, snapshots, export. If you have not, read *Developing a Story with Inkhaven* first; its theology-and-philosophy chapter is the short overview this book is the long, worked expansion of. Where a command or a chord appears here, it is shown doing real work on the essay, not defined in the abstract.

The companion essay was written this way and no other. When you reach its Index Locorum and see every cited passage gathered under its source, you will be looking at the last station of the loop you are about to walk. Turn the page, and let us frame the question.

PART I

Framing the Question

CHAPTER 1

1

Framing the question

Every track begins by telling Inkhaven what kind of work this is, but on this track the telling is load-bearing. Set the genre wrong and the readers apply the wrong standard and generate noise; set it right and they apply the standard the essay was written to. So we begin not with prose but with one line of config, and with the discipline of stating the question sharply enough that the tools can help us answer it.

Set the genre that stops demanding proof

Create the project and set its genre. `philosophy`, `theology`, `theological`, and `religious` share a nearby frame; all of them tell the AI readers to stop asking “is this true?” and start asking “does this *hold together*, and have you met the strongest form of the objection?”:

EDIT · `inkhaven.hjson`

```
genre: "philosophy"
language: "english"
```

The language line matters more than it looks. It drives which public-domain scripture translations the adapters reach for in the next chapter — English pulls the World English Bible; a Russian project would pull the Synodal text and the Kuliev Qur’an without another word from you — and it sets the tongue the finished apparatus is labelled in. State it once, here, and the whole pipeline honours it.

INSIGHT

Declaring the genre recalibrates the interrogators. Point a naïve fact-checker at “the noumenal self is uncreated” and it flags an unsupported empirical claim — a category error, not a finding. With genre: `"philosophy"` set, the same readers grant the frame and press the reasoning instead: where is the hidden premise, what does this presuppose, which objection have you not met? On this track that recalibration is not a nicety. It is the difference between a useful reader and a confused one.

State the question as a thesis you could be wrong about

A research question that cannot fail teaches you nothing. Ours has two live halves — a resemblance and a rupture — and we write it down that way, because the tools in Part III reward a claim with genuine tension in it:

EDIT · the working thesis

Kant's transcendental idealism and Latter-day Saint eternal progression both refuse instantaneous, static salvation and make perfection a *trajectory*. The resemblance is real at two joints — the asymptotic moral self of the second *Critique*, and the uncreated self that grounds freedom. It breaks at three — epistemic posture (postulate versus revelation), asymptote versus arrival (endless approach versus real deification), and the kind of God at the end (a moral postulate versus an embodied, continuous being).

That is not yet an argument; it is a map of where the argument will have to do work. We will spend the rest of the book testing each joint against what the sources actually say — and, more than once, moving a joint when the sources refuse to hold it where we first put it.

ASK YOURSELF

Before you gather a single source, can you name the place your thesis is most likely to be wrong? For us it is “asymptote versus arrival” — Kant may not permit even the endless *approach* we are crediting him with. Write your weakest joint down now. The confront pass in Chapter 5 will go looking for it, and it helps to have named it first.

Sharpen it with the Dialectician

Inkhaven's Research Assistant carries a reader whose whole job is to press a question before you commit to it. `/socrates` points the Inner Socrates Dialectician — the `philosophical-reader` persona — at your corpus and returns the questions your position must answer to stand. Run it early, on the bare thesis, and it will tell you what you have not yet earned:

› /socrates eternal progression and Kant's postulates

The Dialectician asks: (1) You call both views a “trajectory” — but does Kant’s endless approach *to* holiness and the Saints’ progression *into* godhood share a direction, or only a shape? (2) “Refuse instantaneous salvation” — refused by whom, and on what authority: a limit of reason, or a revealed fact? (3) If the gap between creature and God never closes for Kant and does close for the Saints, is your “resemblance” a similarity of doctrine, or only of grammar?

Every one of those questions became a section of the finished essay. That is the point of asking them now: the Dialectician is not answering your question, it is telling you which questions are load-bearing — so you gather to answer *those* and not waste a single source on the parts that were never in doubt.

PITFALL

The characteristic failure of this track is the straw man — defeating a weak version of the view you oppose and mistaking it for a victory. Notice that the Dialectician’s first question already guards against it: it refuses to let “trajectory” mean the same thing for both traditions without proof. Keep the readers pointed at your own argument, not a caricature of the other side. An argument that beats the best opposing case is worth something; one that beats a cartoon persuades no one who matters.

WHAT YOU LEARNED

- Set genre: "philosophy" (or theology) **first** — it recalibrates the readers to stop demanding empirical proof and start pressing coherence, and it is load-bearing on this track, not cosmetic.
- The language line drives which public-domain scripture translations the adapters reach for and the tongue the finished apparatus is labelled in — state it once, here.
- Write the thesis as a **resemblance and a rupture** — a claim with genuine tension — and name the joint most likely to be wrong before you gather anything.
- Run /socrates on the bare thesis to surface the load-bearing questions; gather to answer **those**, and keep the Dialectician pointed at your own argument to forestall the straw man.

PART II

Gathering the Sources

CHAPTER 2

2

Gathering the primary sources

Philosophy and theology are conversations centuries deep, and a position that recalls its sources from memory is weaker for it — and, with a large language model in the loop, actively dangerous, because a model will confabulate a plausible verse or a Kantian page number that says exactly what you wished it said. The discipline of this chapter is the opposite: ingest the *real* texts, keep each with its provenance, and let every later judgement rest on a passage that is actually on disk. Our three primary sources are Kant, the Bible, and the Book of Mormon. All three are public-domain, and all three come in through one command each.

Kant, from the public domain

The Research Assistant (inkhaven research) is where gathering happens. / **gutenberg** searches the keyless Project Gutenberg catalogue and ingests a public-domain book's full text as a research source, chunked and embedded so the relevant passages are retrievable later:

› /gutenberg Kant Critique of Pure Reason

Ingested **The Critique of Pure Reason** (Immanuel Kant, tr. Meiklejohn) as a research source — 214 chunk(s). Cited as @kant - cpr.

Other matches (re-run with an exact title): Kant, Critique of Practical Reason · Kant, Fundamental Principles of the Metaphysic of Morals.

Run it a second time for the second *Critique* — the *Critique of Practical Reason* is where Kant's immortality postulate actually lives, and our thesis leans on it — and Inkhaven files a citation for each to the Sources book automatically as it ingests. We now have Kant on disk, in a translation whose public-domain status is not in doubt, retrievable by meaning rather than by our memory of where a passage sat.

TERM Public-domain sourcing

The rule, on this track, that every ingested source is a text free of copyright restriction — Gutenberg's catalogue, the Internet Archive's scans, and the public-domain scripture translations the adapters reach for. It is a discipline, not a limitation: it keeps the corpus legally clean, and it happens to align with the texts a serious study of a historical tradition wants anyway — the editions the tradition itself was formed on.

The scripture adapters

Scripture is not an ordinary source. A book of it is a structured thing — chapter and verse, surah and ayah — and the unit a theologian cites is not the work but the *passage*. Inkhaven has three adapters built for exactly this shape, each keyless and each drawing only on public-domain translations:

EDIT · the three adapters

`/bible <book> <chapter>` — a Bible passage, verse by verse

`/quran <surah>` — a surah of the Qur'an, ayah by ayah

`/bookofmormon <book> <chapter>` — a Book of Mormon passage, verse by verse

Our essay engages the Bible and the Book of Mormon, so we pull the passages the thesis turns on. The Sermon on the Mount's command to perfection is the Biblical hinge of "eternal progression"; Moroni's promise of being "perfected in Christ" is its Book-of-Mormon counterpart:

› `/bible Matthew 5`

Ingested **The Holy Bible · Matthew 5 (WEB)** — 48 verse(s), 4 chunk(s).

Cited as `@bible` — reference a locus with `@bible[Matthew 5:<verse>]`.

› `/bookofmormon Moroni 10`

Ingested **The Book of Mormon · Moroni 10 (1830)** — 34 verse(s), 3 chunk(s).

Cited as `@book-of-mormon` — reference a locus with `@book-of-mormon[Moroni 10:<verse>]`.

Two details in those replies do quiet, important work. First, the text is *verse-structured*: each chunk carries its verse references alongside the text, so when a passage is retrieved later you see exactly which verses it covers — and the reference is quotable straight into a citation. Second, each work is cited under a *stable key* — `bible`, `book-of-mormon`, `quran` — no matter which chapter seeded it. Ingest Matthew today and John next week, and both file under the one `@bible` key, so every passage you ever cite from scripture gathers under its work in the Index Locorum of Chapter 6.

INSIGHT

The stable cite key is what turns a pile of ingested chapters into a citable edition. Because every Bible passage shares the `bible` key, `@bible[John 3:16]` and `@bible[Matthew 5:48]` are two *loci* of one source, not two unrelated citations — and the machinery that builds the index of cited passages can group them under “The Holy Bible” without your telling it they belong together. You gather chapters; you cite passages; the key holds the two together.

Native tongue, native text

The `language` line you set in Chapter 1 is not decoration here. The adapters choose a public-domain translation *by the project’s language*: English draws the World English Bible and, for the Qur’an, the Sahih International rendering; Russian draws the Synodal Bible and Kuliev’s Qur’an; French, German, and Spanish each draw a public-domain translation of their own. A Russian-language study of the same question would run `/bible Иоанна 3` — the book named in Russian — and receive the Synodal text, in Cyrillic, with no further configuration.

NOTE

You are never locked to the default.

`research.scripture.bible_translation` forces a specific translation code regardless of the project language, and

`research.scripture.quran_translation` does the same for the Qur’an — set it to `quran-uthmani` when you want the Arabic original alongside a translation. The Book of Mormon is English-only by necessity: its 1830 text is public-domain, but the modern translations are under copyright, and this track sources public-domain only.

The same discipline, other doors

Two more adapters round out the primary-source gathering when a tradition’s texts are not scripture in the narrow sense. `/archive` searches the Internet Archive’s

public-domain scans — a way in to older translations, patristic editions, or a philosopher not in Gutenberg — and `/wikisource` pulls a public-domain page from Wikisource *in the book's language*, so a native author reaches native texts. Both ingest exactly as the scripture adapters do — chunked, provenance-tagged, auto-cited — so from the corpus's point of view a surah, a Gutenberg *Critique*, and an Archive scan are the same kind of thing: a retrievable source with a name.

PITFALL

Do not paste a passage you “remember” into your notes and treat it as grounded. The entire point of ingesting the real text is that later stages retrieve from *it* and not from anyone's recollection — yours or the model's. If a passage is not in the corpus, it cannot be retrieved, related, or confronted; it is just an assertion wearing a citation's clothes. Gather the source first; cite from what you gathered.

WHAT YOU LEARNED

- **Gather the real texts, never memory:** `/gutenberg` ingests a public-domain book (Kant's two *Critiques*) and auto-files its citation; a model's recalled verse or page number is confabulation and must not be trusted.
- **The scripture adapters** — `/bible`, `/quran`, `/bookofmormon` — ingest passages verse-by-verse from public-domain translations, each keyless, each auto-cited under a **stable key** so every passage of a work gathers under one source.
- **Language drives translation:** the project's `language` selects the public-domain translation (English→WEB/Sahih, Russian→Synodal/Kuliev, ...); override per work with `research.scripture.bible_translation / quran_translation`.
- **`/archive` and `/wikisource`** extend the same public-domain, provenance-tagged, auto-cited discipline to non-scriptural primary texts — a surah, a Gutenberg *Critique*, and an Archive scan are one kind of thing to the corpus.

CHAPTER 3

3

The corpus and its provenance

Ingesting the sources gave us texts the assistant can retrieve; it did not yet give us *claims*. A study advances on assertions — “Kant grounds freedom outside the causal order”, “the Book of Mormon asserts real deification” — and each of those must be traceable to the passage that warrants it. This short chapter turns the raw sources into a workbench of provenance-tagged facts: the ground the argument will stand on, and the ground the tools in the next chapter will test.

Facts and notes — two shelves, two confidences

Inkhaven keeps two system books for what you learn. The **Facts** book is for claims you are prepared to build on; the **Notes** book is for the speculative, the half-formed, the “check this later.” The Research Assistant files to either with a mandatory confirmation step — nothing enters your ground truth without your seeing it first — and each entry records *where it came from*:

› /fact "Kant postulates immortality as a demand of practical reason, not a theoretical proof" → kant/immortality

Extracted from the last grounded answer.

Source: @kant - cpr (Critique of Practical Reason, retrieved passage).

Confirm insert to **Facts** › **kant** › **immortality**? [y] — Inserted.

The arrow routes the fact to a place in the Facts tree, so the corpus grows an organisation as it grows — a **kant** branch, a **bible** branch, a **book-of-mormon** branch. That organisation is not cosmetic: the contradiction and convergence passes in the next chapter cluster their findings by exactly these branches, so a tree that mirrors your argument gives you a report that mirrors it too.

TERM Provenance

The record, kept per fact, of where a claim came from — which source, retrieved from which passage, at what confidence tier. Provenance is what lets a later reader (human or machine) ask not “is this asserted?” but “on whose authority, and can I go read it?” On this track it is the difference between a claim you can defend before a critic and one you merely remember being persuaded of.

Declare your axioms

Some of what an argument rests on is not a finding but a *commitment* — a position you take as given and ask the reader to grant for the sake of the argument. “We read Kant’s postulates as regulative, not constitutive” is that kind of claim: not true or false in the world, but a stance the essay adopts. Record these in the Facts book

and mark each *undisputed* — select it in the Facts tree and press **u**, and it takes the ✱ glyph. An undisputed fact sits outside the trust ladder: true within your work by decree, with no external truth to test it against.

Marking axioms is not bookkeeping. It changes what the tools do. The **/undisputed** check asks of your declared commitments not “are these true?” but “do these contradict *each other*?” — the mechanical half of the discipline this whole track is about. An argument may be built on any axioms it likes, so long as they cohere; **/undisputed** is the reader that holds you to the “so long as.”

INSIGHT

The distinction between a *fact* (grounded in a source, checkable) and an *axiom* (declared, coherent-or-not) is the distinction the whole track runs on. Facts answer to their provenance; axioms answer only to one another. Keep them on separate footing — grounded claims cite a passage, commitments wear the ✱ — and every later tool knows which standard to apply. Blur them, and the fact-checker will demand proof of a commitment while the coherence check ignores a claim that needed a source.

See where everything came from

At any point, **/sources** lists each fact beside its recorded provenance — the whole corpus as a ledger of who-said-what. It is worth running before the interrogation begins, because a contradiction between two facts means something very different depending on whether they came from *one* source contradicting itself or *two* sources disagreeing — and that distinction, cross-source versus within-source, is one the next chapter’s engine reads straight off the provenance you are looking at here.

WHAT YOU LEARNED

- Turn sources into **claims**: `/fact` files a grounded, provenance-tagged assertion to the **Facts** book (confirmation required); `/note` holds the speculative in the **Notes** book.
- Route facts into a tree that **mirrors your argument** (`→ kant/immortality`) — the contradiction and convergence passes cluster their findings by those same branches.
- Separate **facts** (grounded, checkable) from **axioms** (declared commitments): mark axioms undisputed with `u` (the `⌘`), and `/undisputed` checks them for coherence with one another, not truth.
- `/sources` shows the whole corpus as a provenance ledger — and cross-source versus within-source, which the next chapter's engine reads off that provenance, is the difference between two sources disagreeing and one contradicting itself.

PART III

Interrogating and Reading

CHAPTER 4

4

Interrogating the corpus with SCHOLAR

We have the sources on disk and a corpus of provenance-tagged claims. Now the work proper begins: finding out what the corpus actually *says* when you hold its pieces against one another. This is where a study of two traditions lives or dies — in the places they agree, the places they strain, and the places they flatly contradict — and it is exactly the work retrieval alone cannot do, because retrieval finds what is *similar*, and an argument turns on what is *opposed*. SCHOLAR is the family of commands built to close that gap.

TERM SCHOLAR

Inkhaven's relation engine — the set of Research Assistant commands that judge how claims and sources *stand to one another*, reporting not just contradiction but its confirmation-counterpart too. Where a fact-checker asks “is this claim true?”, SCHOLAR asks “how does this claim relate to that source, and how do these sources relate to each other?” — and grades the answer, from flat contradiction through tension and qualification to agreement.

Relate a claim to the corpus

The workhorse is `/relate`. Give it a claim, and it retrieves the nearest facts and the nearest ingested source passages, then judges *each* against the claim on a graded scale — contradicts, tension, qualifies, agrees, or silent — splitting the verdict into what stands *against* the claim and what *supports* it. Point it at the first joint of our thesis, the asymptotic self:

› `/relate Kant and eternal progression both make perfection an endless approach rather than a state reached`

Relating your claim to the corpus — 2 against · 3 supporting · 1 silent.

✓ **supporting** — `@kant-cpr` (Ak. 5:122): the highest good requires an endless progress toward complete fitness of the will; immortality is postulated for it.

✓ **supporting** — `@bible` (Matthew 5:48): the command “be perfect” set as a standard, not a report of arrival.

⚠ **tension** — `@book-of-mormon` (Moroni 10:32): “perfected in Christ” reads as a reachable completion, not an endless approach.

⚠ **against** — Kant’s holy will is a limit the finite being never becomes; the approach is permanent, which strains “progression toward parity.”

Read what that did. It did not tell us whether our claim is *true* — that is not a question this kind of writing can be held to. It told us where in our own corpus the claim is *supported* and where it is *strained*, and it named the passage for each. The “tension” on Moroni 10:32 is the essay’s second section arriving unbidden: the sources themselves are telling us that “endless approach” holds for Kant and the Sermon on the Mount but strains against the Book of Mormon’s language of arrival.

INSIGHT

/relate is the tool that makes a resemblance *precise*. “Kant and eternal progression rhyme” is a hunch; “they agree at Matthew 5:48 and the second *Critique*, and pull apart at Moroni 10:32” is a thesis with joints you can name and a critic can check. The graded scale is the whole point — a binary “agree/disagree” would have flattened the most interesting finding, which is that the sources *mostly* align and strain at one identifiable seam.

Find where the sources contradict

/relate judges a claim against the corpus. **/contradict** turns the engine on the corpus itself: it scans the Facts book for pairs that cannot both stand, and — reading the provenance from Chapter 3 — reports each as *cross-source* (two sources disagreeing) or *within-source* (one source, or one tradition’s reading, contradicting itself):

› **/contradict**

Found 2 contradiction(s) — 1 between sources, 1 within a source.

[cross-source] · “the creature never becomes holy; the gap is permanent” $\langle @kant - cpr \rangle$ $\Leftrightarrow$ “man may become as God now is” $\langle @book - of - mormon \rangle$ — asymptote versus arrival: the gap Kant fixes, exaltation collapses.

[within-source] · “postulates are not knowledge claims” $\langle @kant - cpr \rangle \Leftrightarrow$ “we may hold immortality with certainty of moral faith” $\langle @kant - cpr \rangle$ — a tension internal to Kant’s own account of how firmly a postulate may be held.

The cross-source contradiction is the spine of our essay’s third movement — the one place the two traditions do not merely strain but genuinely oppose. The within-source one is a gift: it is a tension *inside Kant*, the kind of thing that makes a study more than a comparison, and we would not have gone looking for it. The engine found it because it read Kant’s own claims against each other, not against the Saints.

Find where they converge

Contradiction has a confirmation-counterpart, and SCHOLAR reports it too. **/converge** scans for pairs that *agree*, and marks each convergence as *triangulated* (independent sources arriving at the same claim — the strong kind) or same-source:

› /converge

Found 1 convergence(s) — 1 triangulated.

[triangulated] · “perfection is set as a standard, not reported as reached” `<@bible>` ≈ “the moral law commands a holiness no finite will attains at a moment” `<@kant-cpr>` — the Sermon on the Mount and the second *Critique* arrive independently at perfection-as-trajectory.

A triangulated convergence is worth more than either source alone: when Kant and the Sermon on the Mount reach the same shape of claim by different roads, “perfection is a trajectory” stops being a reading you imposed and becomes one the sources volunteer. That is the first section of the essay, earned.

Question the whole with the Dialectician

We met **/socrates** in Chapter 1, sharpening the bare thesis. Run it again now, over the *gathered* corpus, and it questions what you have actually assembled — the questions the facts must answer to stand:

› /socrates the relation between the two accounts

The Dialectician asks: You have a convergence on “trajectory” and a contradiction on “arrival.” Is arrival a *further* claim the two traditions add to a shared core, or does it retroactively change what “trajectory” meant for each? If the Saints’ trajectory *ends* and Kant’s does not, were they ever the same trajectory — or only the same word?

That is the essay’s hardest paragraph, handed to us as a question. SCHOLAR detects; the Dialectician interrogates. Used together, they map the corpus and then ask whether the map holds.

The persistent report

Every one of those passes rendered its findings into the chat and, by default, into a *persistent report* on disk. `/report` renders the accumulated picture — every contradiction, convergence, and relation you have found — clustered by topic, so the whole interrogation is one document you return to across sessions:

› `/report`

SCHOLAR REPORT — contradictions, convergences, and relations (updated 2026-07-13)

Contradictions (2) — ▶ arrival: asymptote vs. exaltation ... ▶ postulates: Kant’s internal tension ...

Convergences (1) — ▶ trajectory: Bible ≈ Kant, triangulated ...

Relations to claims (1) — ▶ “endless approach”: 3 supporting · 2 against ...

Because the report is keyed to the corpus, it knows when it has gone stale. Ingest a new chapter or edit a fact, and the next `/report` prepends a warning that the facts have changed since the findings were gathered — a prompt to re-run the passes before you trust an old conclusion:

NOTE

The staleness warning is small and load-bearing. A study is written over weeks; the corpus grows the whole time. A conclusion drawn from `/contradict` on Tuesday may no longer hold after Thursday’s new source, and nothing is easier to forget than to re-run the pass. The report carries an order-independent hash of the fact texts, so it can tell you — without your having to remember — that the ground moved under a finding you were about to cite.

WHAT YOU LEARNED

- **SCHOLAR** grades how claims and sources stand to one another — contradiction **and** its confirmation-counterpart — where retrieval alone only finds the similar.
- **/relate** judges a claim against the corpus on a graded scale (contradicts · tension · qualifies · agrees · silent), naming the passage for each — turning a resemblance into a thesis with joints you can name.
- **/contradict** finds pairs that cannot both stand (cross-source vs. within-source); **/converge** finds pairs that agree (triangulated vs. same-source) — the spine and the shared core of the argument, read off the corpus.
- **/socrates** questions the **assembled** corpus (the Dialectician); **/report** renders the accumulated findings clustered by topic and **flags staleness** when the facts move under a conclusion.

CHAPTER 5

5

Reading the draft

Everything in the last chapter happened in the Research Assistant — a separate room from the one you write in. But a study is drafted *while* the corpus is still growing, and the readers that matter most are the ones you can turn on the paragraph in front of you without leaving the editor. This track keeps three, and they ask different questions of the same prose: one confronts the paragraph with your *sources*, one reads it for its *theological weight*, and one reads it for its *validity* — whether the argument, on its own terms, actually follows. Between them they are the second pair of eyes this kind of writing most needs.

Confront the paragraph against the corpus

The first reader is SCHOLAR's **/relate**, brought to the open paragraph. Put the cursor in a paragraph, press **Ctrl+V** then **?**, and Inkhaven retrieves the nearest facts and source passages, judges each against the prose with the graded relation engine, and emits the findings — anchored to the paragraph — into the Output pane. It is the manuscript-side twin of **/relate**: the same judge, pointed at what you are writing.

Suppose we have just drafted the second-section paragraph, the one crediting Kant with an endless approach that never arrives. We press **Ctrl+V ?**, and:

> **Ctrl+V ?** — confront the open paragraph

✗ **confront** · 2 relation(s) to sources in this ¶ — ^B Tab → Output

✗ **against** — @book-of-mormon (Moroni 10:32): “perfected in Christ” reads as a reachable completion; the paragraph’s “never arrives” over-generalises across the traditions it describes.

✓ **supporting** — @kant-cpracr (Ak. 5:122): the endless progress toward fitness of the will confirms the paragraph’s reading of Kant specifically.

The finding marked **against** is the chord doing its job: our paragraph said “the approach never arrives”; the corpus answers that this holds for Kant but that the Book of Mormon, which the paragraph was gesturing at, says the opposite. The prose reached one clause too far, and the source we had already gathered caught it — while the cursor was still in the sentence. An **against** relation arrives as a warning, a **supporting** one as citable information, a silent source is dropped, and each finding re-badges its paragraph in the chapter tree so an unaddressed tension is visible at a glance.

PITFALL

Do not confront a paragraph with no corpus behind it and read the silence as approval. If you have not yet gathered the sources a paragraph depends on, the chord has nothing to retrieve and honestly reports that nothing bears on the prose — which is not the same as the prose being sound. When confront finds nothing, ask whether that is because the claim is safe or because the source is missing.

Read for theological weight — the Inner Theologian

Confront tests a paragraph against what the sources *say*. It cannot tell you that a sentence carries more moral freight than its prose admits, or that a claim about deification has walked past a consequence it never depicted. That reading is the **Inner Theologian's**, and on this track it is the signature reader — the one built for exactly the kind of writing where the stakes are not factual but moral and theological.

Put the cursor in a passage and press **Ctrl+B J**, then **T**. The Inner Theologian reads the passage through the lens of the traditions in play — not against a fact, but for its weight — and its reflection lands in the Thoughts pane, where a longer response belongs. Turn it on the essay's hardest paragraph, the one asserting that for Kant the gap between creature and God stays permanently open:

› Ctrl+B J → T — the Inner Theologian reads the open paragraph

On “*the creature never becomes holy; the gap remains permanently in place*”: this sentence carries a heavier claim than it states. To fix the gap *permanently* is not only an epistemic modesty about Kant's postulates — it is a theological verdict on the worth of the creature, and the paragraph should own it. Consider what the sentence leaves undepicted: a striving with no term is either the creature's dignity (Kant's reading) or its tragedy (the tradition's objection), and the prose currently lets the reader supply whichever they brought. Name the stake. And weigh the word “never” against Matthew 5:48, which you cite two paragraphs earlier as a *command* to perfection — a command implies a capacity the tradition takes seriously and Kant denies. Is your “never” Kant's, or yours?

Read what that did, because it is different from confront. Confront told us a *source* contradicted the paragraph; the Inner Theologian told us the paragraph was *theologically underweight* — that “never” carries a verdict the prose was smuggling rather than arguing, and that the same word sits in tension with a command the essay had already granted. That is not a fact-check. It is the reading a thoughtful editor in the tradition would give, and it made the paragraph's central word the thing to defend rather than assume.

INSIGHT

The Inner Theologian is the reader this track exists for. Every other tool checks the argument's *joints* — where a claim meets a source, where two claims collide. The Inner Theologian reads the argument's *weight*: the consequence left undepicted, the sentence that decides more than it defends, the moral freight a clause carries past what its prose admits. On a track where correctness is off the table, weight is much of what remains — and it is the thing no fact-checker, and no confront pass, can see. Editing a work of theology *is*, in large part, weighing it; this is the reader that weighs.

Two more ways to reach it. **inkhaven theologian scan** reads the whole manuscript for moral and theological weight in one deterministic, zero-AI pass — the fast way to find the paragraphs that deserve a deeper look — and **inkhaven theologian session** opens that deeper look on a passage from the shell. Run the scan across a finished draft and it hands you the list of passages carrying more than they say; open a session on each, and the Inner Theologian reads it as it read the paragraph above.

Read the argument for rigor

Confront reads a paragraph against your sources; the Inner Theologian reads it for its weight. Neither reads it for its *validity* — whether the argument, taken on its own terms, actually follows. That is the **reasoning-rigor reader's** work, and like the Inner Theologian it is deterministic and zero-AI: it does not judge whether you are right, only whether the prose has slipped into one of the recurring shapes of a weak argument. Press **Ctrl+B J**, then **R**, on the open paragraph:

> Ctrl+B J → R — read the open paragraph for rigor

- ✗ **false dichotomy** — “either freedom or determinism”: presented as an exhaustive binary — is there a third alternative you have ruled out without argument?
- ✗ **question-begging** — “obviously”: asserted as self-evident — is the claim argued, or assumed to spare the argument?
- ✗ **non-sequitur** — “therefore”: a conclusion is drawn with no visible warrant in the paragraph — what supports the inference?

It flags five recurring failures — the false dichotomy that hides a third option, the assertion smuggled in as *obviously*, the straw man that answers a weaker opponent, the overgeneralization a single counterexample would break, and the *therefore* with nothing behind it — each by the cue word that gave it away, and each as an advisory, never a verdict. On our essay it caught the sentence that had leaned on “obviously” where the fourth movement in fact required a careful concession; the prose said the argument was settled, and the reader asked whether it was.

INSIGHT

The three readers of this chapter divide the labour of a good editor exactly. Confront asks *is it true to the sources?* The Inner Theologian asks *does it carry its weight?* The rigor reader asks *does it follow?* — and because that last question is purely formal, it can be answered deterministically, in any of the five project languages, with no model in the loop. A Russian argument is read against Russian cues; the advisory comes back in Russian. The reader does not make your argument valid. It tells you, cheaply and in your own tongue, where a reader who wanted to break it would begin.

Reach it across the whole manuscript with `inkhaven rigor scan` — the same reader as a pre-submission pass, `--signal <category>` to focus on one failure, `--strict` to fail a continuous-integration step on any hit. Mute a category your genre uses legitimately (a rhetorical “never”, say) with `rigor.overgeneralization:false` and its kin; the reader is advisory by design, and you decide which of its questions your work is obliged to answer.

Press the argument — the Socratic personas

Alongside the Theologian, the Inner Socrates roster carries two readers built to grant your frame and press your reasoning. `Ctrl+B J`, then the `philosophical-reader` or `theological-reader` persona, and the reader asks where the hidden premise is, what the argument presupposes, which objection you have not met. Because the genre is set (Chapter 1), even these read your prose as an argument to be tested, not a set of facts to be verified.

Use them deliberately against yourself. The characteristic failure of this track is the straw man — defeating a weak version of the view you oppose. Ask the `philosophical-reader` to state the *strongest* form of the objection you are an-

swering, and to find where your argument only defeats a weaker one. On our essay that pressure produced the fourth-movement concession — that Kant’s “trade” of arrival for honesty is not obviously the wrong one — which is what keeps the study a comparison rather than a polemic.

NOTE

Before you act on any of these readings — narrowing a claim confront flagged, rewriting a sentence the Theologian found underweight — take a snapshot (F5), so the draft you are about to change is always recoverable. The next chapter makes snapshots and the prose readers a discipline of their own; here it is enough to know that a reading worth acting on is worth protecting the draft before you do.

WHAT YOU LEARNED

- This track turns three **manuscript readers** on the growing draft, dividing an editor’s labour: Ctrl+V ? **confronts** the paragraph with your sources (is it true to them?), the **Inner Theologian** (Ctrl+B J → T) reads it for **theological weight** (does it carry it?), and the **rigor reader** (Ctrl+B J → R) reads it for **validity** (does it follow?).
- Confront catches a claim that reaches past its evidence; the **Inner Theologian** catches the sentence that carries more moral freight than it admits — the consequence left undepicted, the word that decides more than it defends.
- The **reasoning-rigor reader** is deterministic and zero-AI: it flags five recurring failures — false dichotomy, question-begging, straw man, overgeneralization, non-sequitur — by the cue word that gave each away, in any of the five project languages. Advisory, never a verdict; inkhaven **rigor scan** runs it as a pre-submission pass (--strict for CI).
- Press your own argument with the **philosophical-reader / theological-reader** Socratic personas — ask for the **strongest** objection, to defeat the best opposing case rather than a caricature — and **snapshot** (F5) before acting on any reading.

PART IV

Composing and Producing

CHAPTER 6

6

Composing with loci

A study of scripture and philosophy does not cite *works*; it cites *passages*. “As Kant argues” is worthless to a reader who wants to check you; “as Kant argues at A51/B75” sends them to the exact place. The unit of citation on this track is the locus — the verse, the surah-and-ayah, the Academy-edition page — and Inkhaven cites loci natively, then gathers every one you use into an Index Locorum: the scholarly apparatus that lists, per source, every passage the work engages.

TERM Locus

A specific passage within a source — John 3:16, A51/B75, Ak. 5:122, Moroni 10:32 — as opposed to the source as a whole. In theology, classics, and law the locus is the true unit of citation: the tradition is a conversation about passages, and a claim about a source is checkable only when it names the passage it rests on.

Cite a locus

You already have the cite keys — every ingested source was auto-filed under one, and scripture under a stable key (**bible**, **book-of-mormon**). To cite a locus, write the key and put the passage in brackets:

EDIT · in the manuscript

As the second Critique argues `@kant-cpr[Ak. 5:122]`, holiness is an infinite task; the Sermon on the Mount sets it as a standard, not a report `@bible[Matthew 5:48]`; yet Moroni promises a real completion `@book-of-mormon[Moroni 10:32]`.

That bracketed form is not an Inkhaven invention — it is native Typst, where the bracket after a citation is the *supplement*. Which means the on-disk prose needs no special handling and no transformation: `@bible[Matthew 5:48]` is a real citation to the **bible** source with the supplement “Matthew 5:48”, and it renders and resolves exactly as a page-numbered citation would. Inkhaven’s only addition is to *harvest* these loci for the index.

The cite picker helps you not to mistype a key: `Ctrl+V @` opens a fuzzy search over the Sources book and drops the matching `@key` where the cursor sits. You add the `[locus]` yourself — the picker gets the key right, and the passage is yours to name.

NOTE

Because a locus is a native Typst supplement, the key must resolve to a real bibliography entry — the same requirement any citation carries. `@kant-cpr[Ak.`

`5:122]` compiles only if `kant-cpr` is in your Sources book; a locus on a key with no entry is a compile error, exactly as a citation to a missing reference would be. This is a feature, not a friction: it means the index of passages can never list a source you never actually filed.

Build the Index Locorum

Once the draft cites loci, `inkhaven index-locorum` walks the whole manuscript, harvests every `@key[locus]`, resolves each key to its source title from the Sources book, and renders the apparatus — grouped by source, the loci within each sorted naturally, so `Matthew 5:2` precedes `Matthew 5:48`:

```
inkhaven index-locorum           # Markdown to the
terminal
inkhaven index-locorum --format typst # a Typst chapter
inkhaven index-locorum --format json  # structured, for
tooling
```

The Markdown it produces reads like the back of a scholarly monograph:

EDIT · `index-locorum` output

```
## The Book of Mormon (`@book-of-mormon`)
- Moroni 10:32 — On arrival

## The Critique of Practical Reason (`@kant-cpr`)
- Ak. 5:122 — On trajectory

## The Holy Bible (`@bible`)
- Matthew 5:48 — On trajectory
```

Each locus carries the chapter it was cited in, so the index is not only a list of passages but a small map of where in your argument each one does its work. The

stable scripture keys from Chapter 2 pay off here: every Bible passage you ever cited, from whichever ingested chapter, gathers under the one “The Holy Bible” heading.

Validate the loci, and let them canonicalize

A citation to a passage is only as good as the passage being real. `@bible[John 3:sixteen]` compiles without complaint — Typst has no idea “sixteen” is not a verse — and quietly ships a reference no reader can follow. So a locus is checked against a *reference scheme*: a grammar for what a well-formed reference to that source looks like. The three scripture keys carry built-in schemes (`{book} {ch}:{v}` for the Bible, `{surah}:{ayah}` for the Qur’an), so their loci validate with nothing to configure; for anything else — Kant’s A/B pagination, a Stephanus number — you declare the scheme once:

EDIT · `inkhaven.hjson`

```
sources: {
  ref_schemes: {
    kant-ab: { pattern: "^A\\d+(/B\\d+)?$", format: "A{n}/
B{n}" }
  }
}
```

and name it on the source (a `scheme: kant-ab` line in its Sources entry). A locus that does not match its scheme is flagged — never silently dropped. Catch them from the shell, where `--strict` turns a malformed locus into a failed exit fit for a continuous-integration step:

```
> inkhaven index-locorum --strict
```

```
[?] @bible[John 3:sixteen] — expected {book} {ch}:{v} (The Holy Bible)
[?] @bible[Romans 8] — expected {book} {ch}:{v} (The Holy Bible)
index-locorum: 2 malformed locus(es) → exit 1
```

or catch them as you write: `Ctrl+V c` lints the open paragraph’s loci against their schemes and drops a `[?]` warning into the Output pane for each malformed one,

anchored to the paragraph — the deterministic, editor-side twin of the check above, beside the `Ctrl+V Shift+C` sourcing pass.

INSIGHT

The scheme does a second, quieter service: it *canonicalizes*. Write `@bible[Jn 3:16]` today and `@bible[John 3:16]` tomorrow — or, in a Russian project, `@bible[Иоанна 3:16]` — and the index resolves all three to one entry, “John 3:16”, rather than three near-duplicates scattered under the same source. A reference scheme is not only a gate that rejects the malformed; it is the rule that makes the well-formed *consistent*, so the apparatus reads as one hand wrote it even when your citations, across months of drafting, did not.

Fold it into the finished book

You do not have to run the command by hand. Turn one switch on, and the Index Locorum is assembled into the book automatically, after the bibliography:

EDIT · `inkhaven.hjson`

```
sources: {
  index_locorum: true
}
```

With that set, `inkhaven build` emits an `index_locorum.typ` alongside the `sources.bib` and includes it in the assembled root, right after the reference list — so the finished PDF carries both the bibliography (the works) and the Index Locorum (the passages), each generated from what the manuscript actually cites. It is off by default, because it is a specialised apparatus most books do not want; on this track it is the natural close of the citation story.

INSIGHT

A bibliography answers “which works did this engage?”; an Index Locorum answers “which *passages*?” — and for a study whose entire argument turns on what specific verses and pages say, the second is the more honest record. It is also self-auditing: because it is generated from the loci actually in the prose, it cannot flatter you with a passage you meant to engage but did not. The index is the argument’s contact sheet — every place it touched a source, laid out where a critic can check each one.

WHAT YOU LEARNED

- Cite a **passage**, not a work: `@key[locus]` (`@kant-cpr[Ak. 5:122]`, `@bible[Matthew 5:48]`) is native Typst — the bracket is the supplement — so the prose needs no special handling and the locus renders as a real citation.
- The cite picker (`Ctrl+V @`) gets the **key** right from the Sources book; you name the **locus**. A locus on a key with no entry is a compile error — so the index can never list a source you never filed.
- `inkhaven index-locorum` harvests every `@key[locus]`, resolves titles, and renders the apparatus (Markdown / Typst / JSON) — grouped by source, loci sorted naturally, each tagged with the chapter it was cited in.
- **Validate** loci against a **reference scheme** — built-in for the scripture keys, `sources.ref_schemes` for the rest: `--strict` fails a build on a malformed reference, and `Ctrl+V c` lints the open paragraph as you write. The scheme also **canonicalizes**, so `Jn 3.16`, `John 3:16`, and `Иоанна 3:16` collapse to one entry.
- Set `sources.index-locorum: true` and `inkhaven build` folds the Index Locorum into the finished book after the bibliography — the works **and** the passages, both generated from what the manuscript actually cites.

CHAPTER 7

7

Revising the prose

The readers of Chapter 5 tested the argument — whether a claim met its sources, whether a sentence carried its theological weight. But an argument is delivered in prose, and prose has a second set of demands the interrogators do not touch: it must read well, it must be grammatical, and — because revising means changing what you have already written — it must be *safe to change*. This chapter is the polish pass and the safety net beneath it. Its three tools are older than anything in the SCHOLAR family, and on this track they do the quiet, unglamorous work that lets the argument be heard.

Snapshot before you revise

Revision is destruction with intent: to improve a paragraph you must first be willing to break it, and a study rewritten under a reviewer's questions can go through a dozen versions of its hardest page. Inkhaven's snapshots make that safe. Press **F5** in the editor and the current paragraph is captured — a frozen version you can always return to — before you touch a word:

TRY IT

Put the cursor in the “asymptote versus arrival” paragraph the Inner Theologian flagged in Chapter 5. Press **F5** to snapshot it. Now rewrite it freely — narrow the “never”, name the stake, cite the tension with Matthew 5:48. If the rewrite is worse than the original, press **F6** to open the snapshot picker and restore the version you started from. You have lost nothing by being bold.

F5 creates a snapshot; **F6** opens the picker for the open paragraph, where you can read the earlier versions and restore one; and **Ctrl+F6** opens the project-wide snapshot browser, the whole manuscript's version history in one place. Because a snapshot records *which* draft a paragraph was in, the readers can even attribute a finding to the version it was made against — so you can see that the Theologian's objection was answered by the rewrite, not merely forgotten.

INSIGHT

Snapshots change how boldly you revise. The writer who cannot cheaply undo a change edits timidly — tinkering at the edges of a paragraph they are afraid to break. The writer who can snapshot in one keystroke rewrites the whole thing, reads it, and reverts if it failed. On a track where the hardest paragraphs are rewritten again and again under adversarial reading, that difference — timid tinkering versus fearless revision — is much of the difference in the finished prose.

Grammar-check the passage

With the draft safe to change, check that it is *correct*. Press **F7** on a paragraph and Inkhaven runs a grammar check over its prose. The check keys off the project's **Language**: an English essay is read against English grammar, and a Russian, French, German, or Spanish study is read against its own — the same setting from Chapter 1 that chose the scripture translations now choosing the grammar the prose is held to.

Grammar on this track has one wrinkle worth naming. A study of Kant is thick with technical terms — *noumenal*, *apodictic*, the Latin *a priori* — and with quoted scripture in an older register. The check reads these as prose, so treat its verdict as advice, not law: it will catch the genuine slip — the subject that lost its verb across a long subordinate clause, the number that stopped agreeing — and it will occasionally query a term of art that is exactly right. Take the first; wave off the second.

Read for craft — the Inner Editor

Grammar asks whether a sentence is correct. The **Inner Editor** asks whether it is *good*. It is the prose-craft reader — the one that notices the buried subject, the sentence carrying three clauses that wanted to be two, the paragraph whose best line is its third rather than its first — and it reads your style the way an attentive line-editor would.

Open its overview with **Ctrl+V 0** (O for *Observe*, the editor's defining act), and from there **E** engages the open paragraph; its observations land in the Output pane under the  glyph. Turn it on the essay's dense third movement, and:

› Ctrl+V 0 → E – the Inner Editor reads the open paragraph

 The sentence beginning “The differences are at least as important” opens the movement on an abstraction; the concrete claim — three axes of divergence — arrives only in the next clause. Consider leading with the three.

 “Ontologically continuous with humanity” is doing heavy work in a subordinate position; a reader meets the essay's key term where the sentence is already tired.  Three sentences in this paragraph open with “For Kant”; the parallelism is deliberate, but the third lands as repetition rather than rhythm.

None of that is about whether the argument is *right* — the Theologian and the confront chord already asked that. It is about whether the argument is well *made*: whether the key term arrives where the reader is fresh, whether the rhythm is intended or accidental, whether the concrete claim leads or trails. The Inner Editor is genre-aware, so with genre: "philosophy" set it reads for the virtues this kind of prose wants — clarity under abstraction, a controlled sentence, a claim placed where it can be seen — rather than for a novelist's scene-craft.

PITFALL

The Inner Editor observes; it never rewrites. Nothing it says changes a word of your prose — it emits an observation and leaves the decision to you, exactly as the Theologian and confront do. This is deliberate, and it is the same principle across every reader Inkhaven carries: the tools read, and the author writes. Do not wait for a reader to fix a sentence. It will only ever tell you the sentence could be better; the fix is yours, and so is the credit.

The order of the pass

The three tools have a natural order. Snapshot first (F5), so the revision is safe. Then read for craft (Ctrl+V 0 → E) and correctness (F7), and make the changes the readings earn. If a change goes wrong, revert (F6) and try again. The argument was settled in Chapter 5; this pass settles the prose that carries it — and settles it without fear, because the version you started from is one keystroke away the whole time.

WHAT YOU LEARNED

- **Snapshot before you revise:** F5 captures the open paragraph, F6 opens the picker to restore an earlier version, **Ctrl+F6** browses the whole manuscript's history — so you can rewrite a hard paragraph boldly and revert if the rewrite fails.
- **Grammar-check** with **F7** — it reads against the project's **language**, so the same setting that chose the scripture translations chooses the grammar; take its catches on genuine slips, wave off its queries on terms of art.
- **The Inner Editor** (**Ctrl+V 0 → E**) reads for **craft**, not correctness — the buried subject, the mislaid key term, the rhythm that became repetition — genre-aware, emitting  observations to Output.
- Every reader **observes and never rewrites**: the tools read, the author writes. Snapshot first, read for craft and correctness, act on what the readings earn, and revert fearlessly when a change goes wrong.

CHAPTER 8

8

Producing the essay

The argument has been framed, grounded, interrogated, confronted, and cited. What remains is to turn the tree of paragraphs into something a reader can hold — and to do it so that the finished essay carries, on its own pages, the evidence that it was built the way this book describes: a reference list of the works, an index of the passages, and front matter fit for a journal or an anthology.

The front matter of an essay

An essay is more than its body. Declare its front matter — the abstract, the author, the keywords — in one config block, and it renders into the title page in the book's language:

EDIT · `inkhaven.hjson`

```
frontmatter: {
  abstract: "Kant's transcendental idealism and Latter-day
    Saint eternal
    progression both refuse instantaneous salvation, making
    perfection a
    trajectory. This essay locates two joints where the
    resemblance holds and
    three where it breaks, and argues the family resemblance
    is real but the
    traditions diverge on how confidently one may claim it."
  keywords: ["Kant", "eternal progression", "postulate",
    "deification"]
  authors: [
    { name: "Vladimir Ulogov", affiliation: "Independent",
    corresponding: true }
  ]
}
```

Leave the block empty and nothing renders, so a book that is not an essay is untouched; fill it, and the labels — *Abstract*, *Keywords*, *Corresponding author* — follow `language`, so a Russian study reads in Russian without a second setting.

Render

Then `render.export pdf` produces the essay with its reference list assembled from the Sources you gathered — every citation filed from a real public-domain text — and, with the switch from Chapter 6 set, its Index Locorum after that:

```

inkhaven export pdf --output kant-and-eternal-progression.pdf
inkhaven export epub --output kant-and-eternal-progression.epub
inkhaven export docx --output kant-and-eternal-progression.docx

```

The same tree becomes a PDF for a reader, an EPUB for a shelf, or a DOCX for an editor who redlines in a word processor. Nothing about the argument changes across them; the press does.

NOTE

Two of the earlier tools reappear at production time. The bibliography assembles from the Sources book, so the works you auto-filed while gathering become the reference list with no re-typing; and because every locus in the prose is a real citation to one of those works, a passage cited on a source you never filed would fail the build — the essay cannot ship claiming a source it does not carry.

The essay this produced

This is not a hypothetical pipeline. The companion volume in this directory — *Kant's Transcendental Idealism and Eternal Progression* — is the essay this exact process produced. Its two opening movements are the convergence and the relation `/relate` surfaced in Chapter 4; its third is the cross-source contradiction `/contradict` found; its hardest paragraph is the question `/socrates` asked; and its apparatus is the bibliography and the Index Locorum this chapter renders. Read it beside this one and you can trace every claim back to the station of the loop that produced it.

INSIGHT

The measure of this track is not that Inkhaven wrote the essay — it did not, and must not. The measure is that every claim in the finished essay can be walked back to a passage in a public-domain source, that the tensions were found in your study rather than in a reviewer's, and that the apparatus at the back is generated from what the prose actually cites and not from what you hoped it did. The essay is yours; its *checkability* is what the desk gave you.

WHAT YOU LEARNED

- Declare the essay's **front matter** (abstract, author, keywords) in one config block — it renders into the title page in the book's **language**, and an empty block renders nothing.
- **export pdf | epub | docx** renders the same tree for a reader, a shelf, or an editor — the **bibliography** assembles from the Sources you auto-filed while gathering, with no re-typing.
- Because every locus is a real citation, the build **cannot ship** a passage on a source you never filed — production is the last check, not just the last step.
- The companion essay in this directory is the literal output of this loop — every movement traces to the SCHOLAR pass that produced it, and the back matter to the bibliography and Index Locorum this chapter renders.

PART V

Closing

CHAPTER 9

9

The discipline, in one page

We began with a hunch — that Kant and eternal progression rhyme — and ended with an essay that says exactly where they rhyme and exactly where they do not, every claim walkable back to a public-domain passage. Nothing in that outcome came from a machine that knows theology. It came from a discipline the desk made easy to keep.

What the tools actually did

They did not make the argument. They made it *checkable* — and on a track where correctness is off the table, checkable is the whole of what credibility can mean:

- **Gathering** replaced memory with sources. `/gutenberg` brought Kant; `/bible` and `/bookofmormon` brought scripture, verse by verse, from public-domain translations chosen by the project's language. Nothing was recalled; everything was retrieved.
- **SCHOLAR** replaced a hunch with joints. `/relate` graded the resemblance passage by passage; `/contradict` found the one place the traditions truly oppose, and a tension inside Kant we would not have sought; `/converge` earned the shared core; `/socrates` handed us the hardest question.
- **Confront** (`Ctrl+V ?`) read each paragraph back against the corpus as it was written, catching a claim that reached one clause past its evidence while it was still cheap to fix.
- **The Inner Theologian** (`Ctrl+B J → T`) read each paragraph for its *weight* — the verdict a word was smuggling, the consequence a claim walked past — the reading no fact-check can give and the one this track exists for.
- **The prose readers and the snapshot** did the quiet work: the Inner Editor (`Ctrl+V 0`) read for craft, `F7` for grammar in the book's own language, and `F5` snapshots let the hardest paragraphs be rewritten boldly and reverted without loss.
- **Loci and the Index Locorum** made the citation the passage, not the work — and generated, from the prose itself, the apparatus a serious study is judged by.

INSIGHT

The through-line of every stage is the same refusal: the refusal to let a claim rest on anyone's recollection — the author's or the model's. Sources are ingested, not remembered; relations are judged from retrieved passages, not asserted; the index is harvested from real citations, not compiled from good intentions. That refusal is the discipline. The tools exist to make keeping it the path of least resistance.

Where to go next

This book worked one question with one set of sources. The same loop scales in three directions. *More traditions*: `/quran` brings the Qur'an under the same stable-key, verse-structured discipline, and `/archive` and `/wikisource` reach public-domain texts beyond scripture. *More languages*: set `language` to Russian, French, German, or Spanish and the adapters draw that tongue's public-domain

translations, the SCHOLAR reports and the Index Locorum heading follow suit, and the whole pipeline runs in the language the study is written in. *More rigour*: the adversarial verdict personas and the coherence checks from the scientific and nonfiction tracks are yours whenever an argument wants a hostile reader.

ASK YOURSELF

What is the one claim in your own project you have been carrying on memory alone — the verse you are *sure* says what you need, the page you *remember* Kant conceding? That is the claim to gather first. Ingest the source, /relate your claim to it, and find out whether the passage agrees with you or only your recollection did. That single act — turning one remembered claim into a checked one — is this entire book in miniature.

Turn now to the companion volume, and read the essay this discipline produced. Every place it names a passage, you will know the passage was on disk before the sentence was written — and that, on this track, is the whole of the achievement.

AFTERWORD

About the author



Vladimir Ulogov.

Vladimir Ulogov has spent decades building infrastructure for distributed systems — the kind of software that watches other software. Early in his career he worked on monitoring and telemetry platforms; later years took him into federated observability, telemetry buses, and the architecture of systems that have to make sense of millions of data points without losing the thread.

Observability, in the end, is a discipline of **coherence** — of never reporting a state the system cannot account for, of insisting that every signal follow from something real. It is not an accident that a tool he built for writers carries the same instinct into a work of theology, where the whole art is to advance nothing the argument cannot answer for.

What makes him slightly unusual in his corner of the industry is a tendency to write his own tools — not in the sense of small utilities, but in the sense of programming languages. The Bund language (its compiler, its VM, its document store, its parser) lives in a long series of Rust crates on crates.io. `rust_dynamic`, `rust_multistack`, `rust_multistackvm`, `bundcore`, `zbus_universal_data_gateway` — each is a building block that exists because the off-the-shelf options didn’t fit the shape of the work. Inkhaven’s research desk, with its provenance ledger and its refusal to let a claim rest on memory, is that same instinct pointed at the humanities.

Why a tool for the unmeasurable

Inkhaven was built to serve every kind of book, and this one — the theological essay, the philosophical argument — is the hardest case for a tool that began in a world of metrics and telemetry. There is no dataset here, no fact to check against the world, no ground truth a monitor could alarm on. And yet the underlying discipline turned out to be the same: hold a large structure of commitments without letting it quietly contradict itself, and never assert what you cannot trace. Observability calls that coherence; philosophy has called it the same thing for rather longer. Inkhaven’s theology-and-philosophy track is the attempt to bring an engineer’s instinct for traceability to a kind of writing that measurement cannot touch — not to settle its questions, which no tool can, but to make its arguments *checkable* in the only sense they admit.

A work of love

Inkhaven is open source. The licence is permissive — you can read it, fork it, study it, modify it, and pass it on. Strictly speaking, the licence also lets you sell it. The author would, gently but firmly, disagree with you doing that. Inkhaven was not designed as a *for sale* project. It is a work of love made for the authors who can least afford to pay for software, and turning it into a commercial product would betray the reason it exists. So please — don’t.

It carries no analytics, no telemetry, no upsell. The binary will never phone home; the project will never have an “Enterprise” tier you have to escape from.

This was a deliberate choice. There are excellent commercial tools for writing. They cost money — which is fine for many writers, and a barrier for many more. Inkhaven exists for the second group: for the graduate student writing a thesis

in divinity on a battered laptop, for the independent scholar working outside any institution's subscriptions, for the pastor drafting in the same terminal where they keep their notes — for anyone who would benefit from a tool that respects their work without asking for a credit-card number.

A note on cooperation

Vladimir believes firmly in the human capacity for mutual help — that we make better work, and live better lives, when we share what we know and what we build. Open source is one of the most concrete expressions of cooperation our era has produced: code read, improved, and passed forward without payment, without permission, by people who will never meet.

If Inkhaven helps you finish your book — whatever kind of book it is — that is enough. If it gives you a chord pattern you adapt into your own tool, that's a gift back to the larger project of making software writers can love. As a society, we achieve the greatest things when we help each other rather than compete with each other.

Where to find more

GitHub	@vulogov — the source for Inkhaven, Bund, and the dozen-plus Rust crates that carry the infrastructure. Issues and PRs welcome.
LinkedIn	/in/vladimirulogov — posts on observability, the occasional long-form essay.
YouTube	@vulogov — talks and walkthroughs from the conference trail.
X / Twitter	@vladimir_ulogov

Whatever you are arguing — if the tool ever gets in the way of the argument instead of out of it, open an issue on GitHub. The author reads them.

THEOLOGY AND PHILOSOPHY WITH INKHAVEN

END OF THE BOOK · INKHAVEN 1.6.20